



Sri Lanka Institute of Information Technology

Financial Risk Prediction

Project Report

Fundamental of Data Mining – IT3051

Group 16/Data Miners

Submitted by:

1. Bandara K M W G L A – IT22089304
2. Jayasinghe J M H C – IT22597342
3. Abeysinghe E A – IT22166692
4. Wicramasinghe D A T N – IT22113504
5. Mesandu W M S – IT22108890

Submitted to:

Dr. Prasanna S. Haddela

Name of the supervisor/lecturer in charge

10/03/2024

Date of submission

Video link : [WhatsApp Video 2024-10-03 at 23.49.40.mp4](#)

Deployed link:

Github link: <https://github.com/IT22089304/Financial-risk-prediction.git>

Contents

| | |
|--|----|
| Introduction | 2 |
| Project Background | 2 |
| Problem Statement | 2 |
| Current Challenge..... | 2 |
| Significance..... | 2 |
| Proposed Solution..... | 3 |
| Impact | 3 |
| Dataset analysis and preparation | 4 |
| Implementation | 6 |
| Data Preparation | 6 |
| Data Preprocessing | 6 |
| Model selection..... | 15 |
| 1.Logistic Regression Model | 16 |
| 2. Random Forest Classifier | 16 |
| 3. Gradient Boosting Classifier..... | 16 |
| 4. K-Nearest Neighbors Classifier..... | 17 |
| 5. Support Vector Machine (SVMs) | 17 |
| Model training | 18 |
| Hyperparameter tuning..... | 19 |
| Cross-validation..... | 19 |
| Frontend | 20 |
| Conclusion..... | 22 |
| Work Distribution..... | 23 |

Introduction

Project Background

This project focuses on building a machine learning model to evaluate financial risks for individuals applying for loans, aiming to reduce defaults and improve credit risk prediction accuracy. In the financial industry, assessing an individual's risk profile is crucial for minimizing loan defaults and ensuring responsible lending practices.

Our model will be developed using a publicly available dataset, which includes data on demographics, financial status, and past credit behavior of individuals. By analyzing this data, the model will predict the individual's financial risk level categorized as Low or High. The ultimate goal is to provide more accurate risk evaluations, helping financial institutions make informed decisions and manage credit risks effectively.

Problem Statement

Current Challenge

Financial institutions face difficulties in accurately assessing the credit risk of individuals applying for loans. Rigid scoring systems and other traditional credit evaluation techniques frequently fall short of capturing the complexity of an applicant's financial history and behavior. With regard to this, it may be challenging to estimate the actual risk of loan default using those approaches, which could result in inaccurate creditworthiness ratings. This challenge leads to higher default rates, increased financial risk for lenders, and inaccurate loan decisions.

Significance

Inaccurate risk predictions have adverse effects on lenders as well as loan applicants. For lenders, surmising incorrectly that an applicant is a lower risk leads to increased defaults on loans and hence lower profitability. Applicants, on the other hand, may be unfairly denied loans or offered unfavorable loan terms based on incorrect risk classification. This affects not only individuals' ability to access credit but also the fairness and efficiency of the overall lending process.

Proposed Solution

The purpose of this project is to create a model based on machine learning that will examine the financial risk status of loan applicants. The model utilizes a dataset comprising demographic data, economic data and credit history data predicting the risk level of the applicant low or high. By leveraging machine learning, the model will help to detect trends and relationships in the data that may not be obtained using the conventional approaches further making the model superior.

Impact

The implementation of a more accurate and data-driven financial risk assessment model will significantly reduce loan default rates, thereby enhancing the safety and profitability of financial institutions. It will also make credit evaluations more equitable by de-biasing the processes creating barriers to high-risk borrowers who have been misclassified in conventional screening methods. In the long term, this project will enhance the practices of lending and borrowing in the whole financial market to be more efficient, accountable and more transparent.

Dataset analysis and preparation

| Variable | Variable Name | Description | Variable Type |
|----------|----------------------|--|---------------|
| 1 | Age | The applicant's age. | numerical |
| 2 | Gender | Gender ("Male", "Female", "Non-binary") | categorical |
| 3 | Education | Education status ("PhD", "Master's", "Bachelor's", "High School") | categorical |
| 4 | Marital Status | Education status ("divorced", "married", "single", "widowed") | categorical |
| 5 | Income | The annual income of the applicant. | numerical |
| 6 | Credit Score | The applicant's credit score, which reflects their creditworthiness. | numerical |
| 7 | Loan Amount | The amount of the loan being applied for. | numerical |
| 8 | Loan Purpose | Type of Loan purpose ("Business", "Home", "Personal", "Auto") | categorical |
| 9 | Employment Status | Employment Status ("Employed", "Unemployed", "Self-employed") | categorical |
| 10 | Years at Current Job | The number of years the applicant has been at their current job. | numerical |
| 12 | Debt-to-Income Ratio | The ratio of the applicant's debt compared to their income. | numerical |
| 13 | Assets Value | The total value of the applicant's assets. | numerical |

| | | | |
|----|-----------------------|--|-------------|
| 14 | Number of Dependents | The number of dependents the applicant is financially responsible for. | numerical |
| 15 | City | The city where the applicant resides. (Too many unique values to list here, some examples - "Port Elizabeth", "South Scott", "Robin haven", "New Heather") | categorical |
| 16 | State | The state where the applicant resides. (some examples-AS, OH, OK, PR) | categorical |
| 17 | Country | The country where the applicant resides. (some examples-Cyprus, Turkmenistan, Luxembourg, Uganda, Namibia) | categorical |
| 18 | Previous Defaults | The number of times the applicant has defaulted on a loan in the past. | numerical |
| 19 | Marital Status Change | Indicates any recent change in marital status. | numerical |
| 20 | Risk Rating | The final risk assessment of the applicant's financial situation (Low or High). | categorical |

Implementation

Data Preparation

The dataset is partitioned into:

- Training set - 80%
- Testing set - 20%

Data Preprocessing

Import the data set

```
[3]: df = pd.read_csv("financial_risk_assessment.csv")
[4]: df.head()
```

| | Age | Gender | Education Level | Marital Status | Income | Credit Score | Loan Amount | Loan Purpose | Employment Status | Years at Current Job | Payment History | Debt-to-Income Ratio | Assets Value | Number of Dependents | City | State | Country |
|---|-----|------------|-----------------|----------------|---------|--------------|-------------|--------------|-------------------|----------------------|-----------------|----------------------|--------------|----------------------|----------------|-------|--------------------|
| 0 | 49 | Male | PhD | Divorced | 72799.0 | 688.0 | 45713.0 | Business | Unemployed | 19 | Poor | 0.154313 | 120228.0 | 0.0 | Port Elizabeth | AS | Cy |
| 1 | 25 | Non-binary | Bachelor's | Widowed | 49427.0 | 766.0 | 36528.0 | Personal | Unemployed | 10 | Fair | 0.143242 | 287140.0 | NaN | New Heather | IL | Nar |
| 2 | 31 | Non-binary | Master's | Widowed | 45280.0 | 672.0 | 6553.0 | Personal | Self-employed | 1 | Good | 0.378890 | NaN | NaN | West Lindaview | MD | Bc Is (Bouve |
| 3 | 18 | Male | Bachelor's | Widowed | 93678.0 | NaN | NaN | Business | Unemployed | 10 | Poor | 0.396636 | 246597.0 | 1.0 | Melissahaven | MA | Honc |
| 4 | 32 | Non-binary | Bachelor's | Widowed | 20205.0 | 710.0 | NaN | Auto | Unemployed | 4 | Fair | 0.335965 | 227599.0 | 0.0 | North Beverly | DC | Pit Isl |

Check duplicate values, missing values & garbage values.

```
[7]: ## Check duplicates
df.duplicated().sum()

[7]: 0

[8]: # Check null count
df.isnull().sum()

[8]: Age                0
Gender                0
Education Level       0
Marital Status        0
Income              1573
Credit Score         1555
Loan Amount          1600
Loan Purpose          0
Employment Status     0
Years at Current Job  0
Payment History       0
Debt-to-Income Ratio  0
Assets Value          1609
Number of Dependents  1571
City                  0
State                 0
Country               0
Previous Defaults     1533
Marital Status Change 0
Risk Rating           0
dtype: int64
```

```
[9]: # Check null count as percentage
```

```
col_num=0
TotalObjects =df.shape[0]
print ("Column\t\t\t\t\t Null Values%")
for x in df:
    nullCount =df[x].isnull().sum();
    nullPercent = nullCount*100 / (TotalObjects)
    print(str(x)+"\t\t\t\t\t "+str(nullPercent))
```

| Column | Null Values% |
|-----------------------|--------------------|
| Age | 0.0 |
| Gender | 0.0 |
| Education Level | 0.0 |
| Marital Status | 0.0 |
| Income | 14.980952380952381 |
| Credit Score | 14.80952380952381 |
| Loan Amount | 15.238095238095237 |
| Loan Purpose | 0.0 |
| Employment Status | 0.0 |
| Years at Current Job | 0.0 |
| Payment History | 0.0 |
| Debt-to-Income Ratio | 0.0 |
| Assets Value | 15.323809523809524 |
| Number of Dependents | 14.961904761904762 |
| City | 0.0 |
| State | 0.0 |
| Country | 0.0 |
| Previous Defaults | 14.6 |
| Marital Status Change | 0.0 |

```
[11]: #identifying garbage values
```

```
for i in df.select_dtypes(include="object").columns:
    print(df[i].value_counts())
    print("*****10")
```

```
Gender
Non-binary    3565
Female        3499
Male          3436
Name: count, dtype: int64
*****
Education Level
Bachelor's    2677
High School   2627
PhD           2624
Master's      2572
Name: count, dtype: int64
*****
Marital Status
Widowed       2713
```

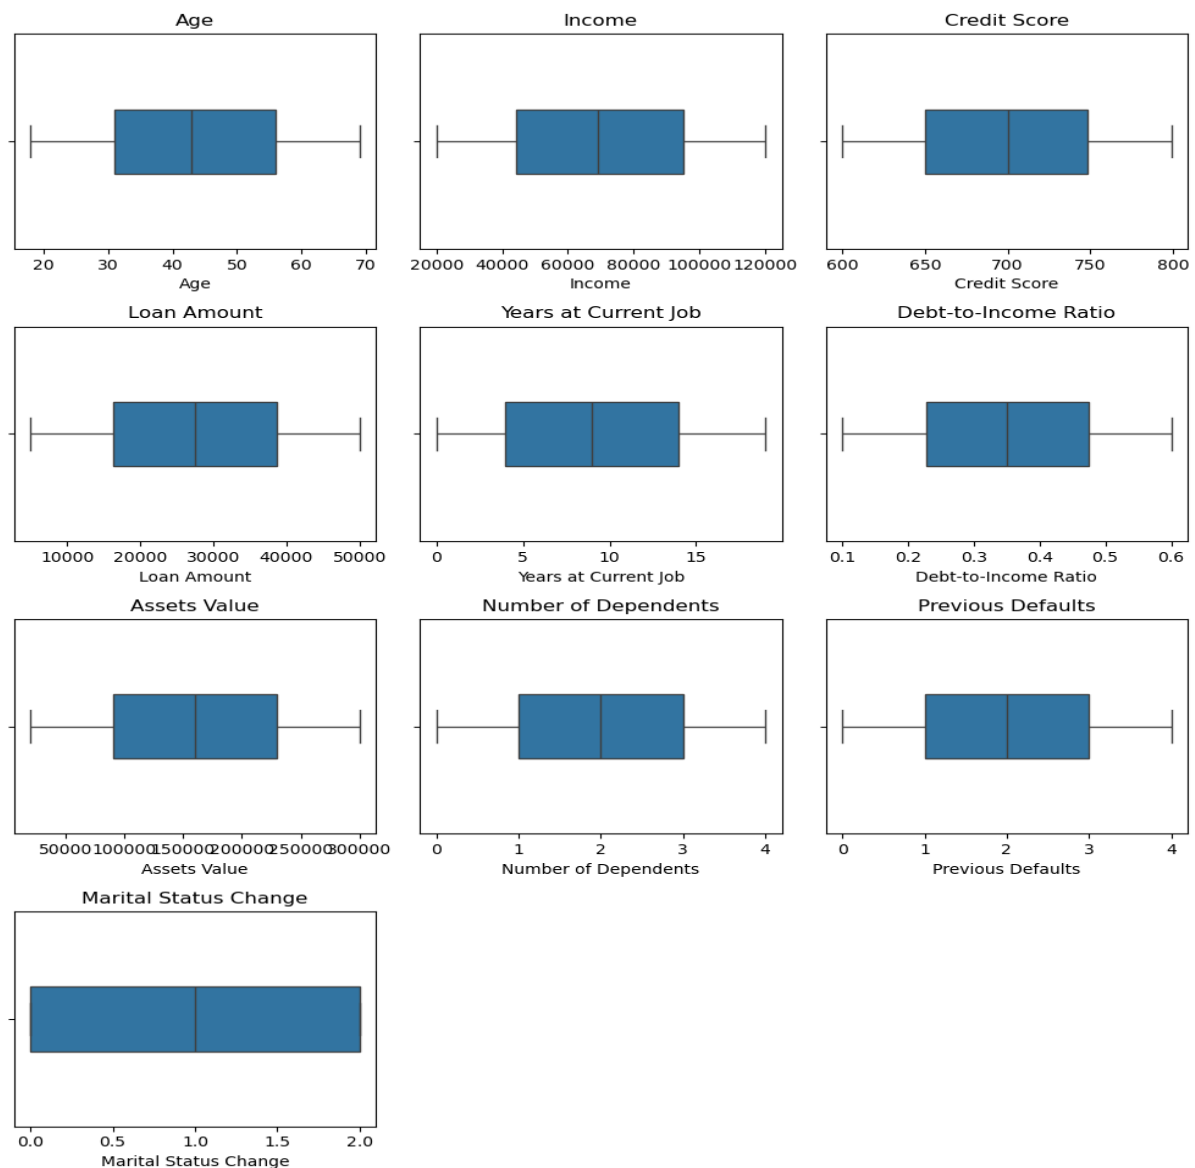
Check data types

```
[10]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10500 entries, 0 to 10499
Data columns (total 20 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   Age                  10500 non-null  int64
1   Gender               10500 non-null  object
2   Education Level      10500 non-null  object
3   Marital Status       10500 non-null  object
4   Income               8927 non-null   float64
5   Credit Score         8945 non-null   float64
6   Loan Amount          8900 non-null   float64
7   Loan Purpose         10500 non-null  object
8   Employment Status    10500 non-null  object
9   Years at Current Job 10500 non-null  int64
10  Payment History      10500 non-null  object
11  Debt-to-Income Ratio 10500 non-null  float64
12  Assets Value         8891 non-null   float64
13  Number of Dependents 8929 non-null   float64
14  City                 10500 non-null  object
15  State                10500 non-null  object
16  Country              10500 non-null  object
17  Previous Defaults     8967 non-null   float64
18  Marital Status Change 10500 non-null  int64
19  Risk Rating          10500 non-null  object
dtypes: float64(7), int64(3), object(10)
memory usage: 1.6+ MB
```


Check outliers

```
[31]: import warnings
import matplotlib.pyplot as plt
import seaborn as sns
warnings.filterwarnings("ignore")
# Select numerical columns
numerical_cols = df.select_dtypes(include="number").columns
# Set up the figure and axes with a smaller size
n = len(numerical_cols)
cols = 3 # Number of columns in the grid
rows = (n // cols) + (n % cols > 0) # Calculate number of rows needed
fig, axes = plt.subplots(rows, cols, figsize=(10, 3 * rows)) # Smaller figure size
axes = axes.flatten() # Flatten the axes array for easy indexing
# Plot each boxplot
for i, col in enumerate(numerical_cols):
    sns.boxplot(data=df, x=col, ax=axes[i], width=0.3) # Adjust width for smaller plots
    axes[i].set_title(col)
# Hide any unused subplots
for j in range(i + 1, len(axes)):
    fig.delaxes(axes[j])
plt.tight_layout()
plt.show()
```



Handle missing values

```
[15]: #Missing value treatments using mean for continous value columns
      for i in ["Income","Credit Score","Loan Amount","Assets Value"]:
          df[i].fillna(df[i].mean(),inplace=True)

[16]: #Missing value treatments using mode for discrete value columns

      from sklearn.impute import SimpleImputer

      # Create an imputer object with most_frequent strategy
      imputer = SimpleImputer(strategy='most_frequent')

      # Fit and transform the data
      df['Number of Dependents'] = imputer.fit_transform(df[['Number of Dependents']])
      df['Previous Defaults'] = imputer.fit_transform(df[['Previous Defaults']])

[17]: df.isnull().sum()

[17]: Age                0
      Gender             0
      Education Level    0
      Marital Status     0
      Income             0
      Credit Score       0
      Loan Amount        0
      Loan Purpose       0
      Employment Status  0
      Years at Current Job 0
      Payment History    0
      Debt-to-Income Ratio 0
      Assets Value       0
      Number of Dependents 0
      City              0
      State             0
      Country           0
      Previous Defaults   0
      Marital Status Change 0
      Risk Rating        0
      dtype: int64
```

Find correlation between Risk Rating column with other categorical columns

```
[20]: from scipy.stats import chi2_contingency
      def cramers_v(contingency_table):
          chi2, p, dof, expected = chi2_contingency(contingency_table)
          n = contingency_table.sum().sum() # Total sample size
          r, k = contingency_table.shape
          cramers_v = np.sqrt(chi2 / (n * (min(r, k) - 1)))
          return cramers_v

      contingency_table = pd.crosstab(df['Gender'], df['Risk Rating'])
      cramers_v_value = cramers_v(contingency_table)
      print(f"Cramér's V for Gender and Risk Rating: {cramers_v_value}")

      contingency_table = pd.crosstab(df['Education Level'], df['Risk Rating'])
      cramers_v_value = cramers_v(contingency_table)
      print(f"Cramér's V for Education Level and Risk Rating: {cramers_v_value}")

      contingency_table = pd.crosstab(df['Marital Status'], df['Risk Rating'])
      cramers_v_value = cramers_v(contingency_table)
      print(f"Cramér's V for Marital Status and Risk Rating: {cramers_v_value}")

      contingency_table = pd.crosstab(df['Loan Purpose'], df['Risk Rating'])
      cramers_v_value = cramers_v(contingency_table)
      print(f"Cramér's V for Loan Purpose and Risk Rating: {cramers_v_value}")
```

```

contingency_table = pd.crosstab(df['Employment Status'], df['Risk Rating'])
cramers_v_value = cramers_v(contingency_table)
print(f"Cramér's V for Employment Status and Risk Rating: {cramers_v_value}")

contingency_table = pd.crosstab(df['Payment History'], df['Risk Rating'])
cramers_v_value = cramers_v(contingency_table)
print(f"Cramér's V for Payment History and Risk Rating: {cramers_v_value}")

contingency_table = pd.crosstab(df['City'], df['Risk Rating'])
cramers_v_value = cramers_v(contingency_table)
print(f"Cramér's V for City and Risk Rating: {cramers_v_value}")

contingency_table = pd.crosstab(df['State'], df['Risk Rating'])
cramers_v_value = cramers_v(contingency_table)
print(f"Cramér's V for State and Risk Rating: {cramers_v_value}")

contingency_table = pd.crosstab(df['Country'], df['Risk Rating'])
cramers_v_value = cramers_v(contingency_table)
print(f"Cramér's V for Country and Risk Rating: {cramers_v_value}")

```

```

Cramér's V for Gender and Risk Rating: 0.024397047231932888
Cramér's V for Education Level and Risk Rating: 0.012680084204131905
Cramér's V for Marital Status and Risk Rating: 0.022375185950592878
Cramér's V for Loan Purpose and Risk Rating: 0.025668056146136632
Cramér's V for Employment Status and Risk Rating: 0.015433015305277563
Cramér's V for Payment History and Risk Rating: 0.019797386112888005
Cramér's V for City and Risk Rating: 0.8802450179057865
Cramér's V for State and Risk Rating: 0.07595302023802004
Cramér's V for Country and Risk Rating: 0.15003672741293336

```

Drop columns

```

[24]: # Drop unnecessary categorical columns
df = df.drop(['State', 'Country', 'Education Level', 'Marital Status', 'Marital Status Change', 'Gender', 'Years at Current Job'], axis=1)

print(df.columns)

Index(['Age', 'Income', 'Loan Amount', 'Loan Purpose', 'Employment Status',
       'Payment History', 'Assets Value', 'City', 'Previous Defaults',
       'Risk Rating'],
      dtype='object')

```

For Ordinal Categorical Data Encoding using Label Encoding

```

[24]: from sklearn.preprocessing import LabelEncoder

label_encoder = LabelEncoder()

# defined order for ordinal columns
#education_order = ["High School", "Bachelor's", "Master's", 'PhD']
payment_history_order = ['Poor', 'Fair', 'Good', 'Excellent']
risk_rating_order = ['Low', 'Medium', 'High']

#df['Education Level'] = pd.Categorical(df['Education Level'], categories=education_order, ordered=True)
df['Payment History'] = pd.Categorical(df['Payment History'], categories=payment_history_order, ordered=True)
df['Risk Rating'] = pd.Categorical(df['Risk Rating'], categories=risk_rating_order, ordered=True)

#df['Education Level Encoded'] = label_encoder.fit_transform(df['Education Level'])
df['Payment History Encoded'] = label_encoder.fit_transform(df['Payment History'])
df['Risk Rating Encoded'] = label_encoder.fit_transform(df['Risk Rating'])

```

Group city by risk rating mean

```
[25]: # Group by City and calculate the mean Risk Rating for each city
city_risk_means = df.groupby('City')['Risk Rating Encoded'].mean()

# Assign groups based on risk Level
df['City_grouped'] = df['City'].apply(lambda x: 'High_Risk_Cities' if city_risk_means[x] > 0.5
                                     else 'Low_Risk_Cities')

print(df[['City', 'City_grouped']])
```

| | City | City_grouped |
|-------|---------------------|------------------|
| 0 | Port Elizabeth | High_Risk_Cities |
| 1 | New Heather | High_Risk_Cities |
| 2 | West Lindaview | High_Risk_Cities |
| 3 | Melissahaven | High_Risk_Cities |
| 4 | North Beverly | High_Risk_Cities |
| ... | ... | ... |
| 10495 | Curtismouth | High_Risk_Cities |
| 10496 | Susanstad | High_Risk_Cities |
| 10497 | South Morganchester | Low_Risk_Cities |
| 10498 | Port Wayne | Low_Risk_Cities |
| 10499 | South Stacy | High_Risk_Cities |

[10500 rows x 2 columns]

```
[26]: df['City_grouped'].value_counts()
```

```
[26]: City_grouped
High_Risk_Cities    8916
Low_Risk_Cities     1584
Name: count, dtype: int64
```

Correlation Matrix

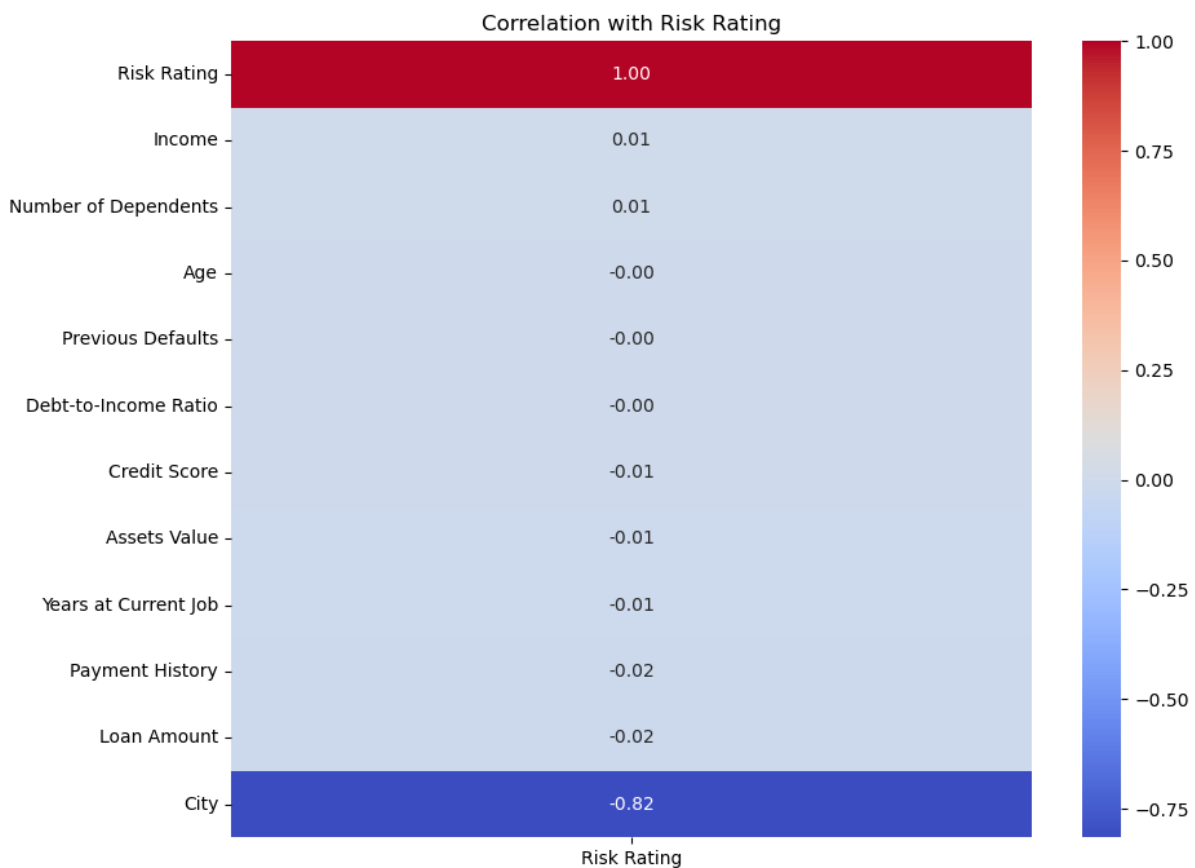
```
[30]: # Select only numerical columns from the dataframe
numerical_cols = df.select_dtypes(include=['int64', 'int32', 'float64'])

# Calculate correlation matrix for numerical columns
corr_matrix = numerical_cols.corr()

# Check if 'Risk Rating' is an existing column
if 'Risk Rating' in corr_matrix.columns:
    # Get the correlation between 'Risk Rating' and other numerical features
    risk_rating_correlation = corr_matrix['Risk Rating'].sort_values(ascending=False)

    # Visualize the correlation matrix as a heatmap using seaborn
    import seaborn as sns
    import matplotlib.pyplot as plt

    plt.figure(figsize=(10, 8))
    sns.heatmap(risk_rating_correlation.to_frame(), annot=True, cmap='coolwarm', fmt='.2f')
    plt.title('Correlation with Risk Rating')
    plt.show()
else:
    print("'Risk Rating' column is not found in the correlation matrix.")
```



Nominal Categorical data encoding with One-Hot Encoding

```
[31]: # List of columns to One-Hot Encode
categorical_columns = [ 'Loan Purpose', 'Employment Status' ]

# Apply One-Hot Encoding to the specified categorical columns
df_encoded = pd.get_dummies(df, columns=categorical_columns, drop_first=False)

# Convert only the newly created one-hot encoded columns to int
one_hot_encoded_columns = df_encoded.columns.difference(df.columns)

# Apply astype(int) only to these new one-hot encoded columns
df_encoded[one_hot_encoded_columns] = df_encoded[one_hot_encoded_columns].astype(int)

# Replace original DataFrame with the encoded version
df = df_encoded.copy()

print("DataFrame after One-Hot Encoding with original columns replaced:")
print(df)
```

Scaling numerical features

```
[32]: from sklearn.preprocessing import MinMaxScaler

# Only scale certain numerical columns
numerical_columns_to_scale = ['Age', 'Income', 'Loan Amount', 'Credit Score', 'Assets Value', 'Debt-to-Income Ratio', 'Years at
df_numerical = df[numerical_columns_to_scale]

# Apply Min-Max Scaling
scaler = MinMaxScaler()
df_scaled = pd.DataFrame(scaler.fit_transform(df_numerical), columns=df_numerical.columns)

# Combine with the rest of the data that doesn't require scaling
df_rest = df.drop(columns=numerical_columns_to_scale)
df = pd.concat([df_rest, df_scaled], axis=1)

print("Final DataFrame after selective scaling:")
print(df.head())
```

```
[33]: # Calculate the correlation matrix
correlation_matrix = df.corr()

# Get the correlation between 'Risk Rating' and other features
risk_rating_correlation = correlation_matrix['Risk Rating'].sort_values(ascending=False)

# Display the correlation values
print(risk_rating_correlation)

Risk Rating                1.000000
Loan Purpose_Business      0.022970
Employment Status_Unemployed 0.014693
Income                    0.006749
Number of Dependents       0.006014
Loan Purpose_Home          0.001526
Age                       -0.002167
Previous Defaults          -0.002766
Debt-to-Income Ratio      -0.003034
Employment Status_Self-employed -0.003292
Credit Score              -0.005079
Loan Purpose_Personal      -0.007322
Assets Value               -0.007482
Employment Status_Employed -0.011398
Years at Current Job       -0.012975
Payment History            -0.015373
Loan Amount                -0.015577
Loan Purpose_Auto          -0.017108
City                      -0.815725
Name: Risk Rating, dtype: float64
```

Split the data

```
[36]: from sklearn.model_selection import train_test_split

# Splitting the data into features (X) and target (y)
X = df.drop(columns='Risk Rating')
y = df['Risk Rating']

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)
```

```
[37]: print(f"Training samples: {X_train.shape}")
print(f"Validation samples: {X_test.shape}")
print(f"Training samples: {y_train.shape}")
print(f"Validation samples: {y_test.shape}")

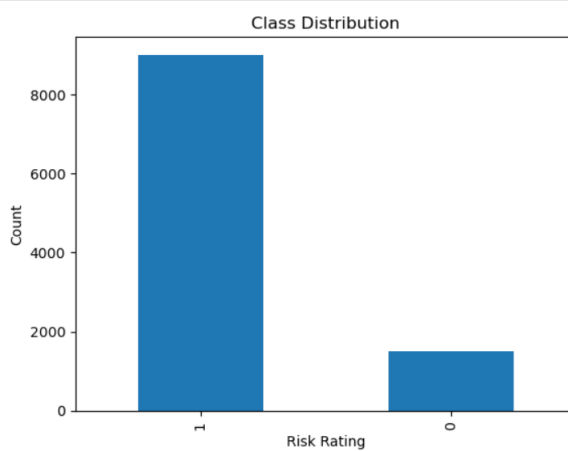
Training samples: (8400, 18)
Validation samples: (2100, 18)
Training samples: (8400,)
Validation samples: (2100,)
```

Handling imbalanced dataset

```
[34]: # Check class column distribution
import matplotlib.pyplot as plt

# Assuming 'df' is your DataFrame and 'Risk Rating' is your target variable
risk_rating_counts = df['Risk Rating'].value_counts()

# Plot the class distribution as a bar chart
risk_rating_counts.plot(kind='bar')
plt.title("Class Distribution")
plt.xlabel("Risk Rating")
plt.ylabel("Count")
plt.show()
```



```
[51]: from sklearn.utils import resample
from sklearn.model_selection import train_test_split

# 'Risk Rating' is the target column
X = df.drop(columns=['Risk Rating']) # Features
y = df['Risk Rating'] # Target

# Combine X and y into one DataFrame for easy manipulation
df_combined = pd.concat([X, y], axis=1)

# Find the value counts for each class in the target column
class_counts = y.value_counts()

# Identify the majority and minority classes
majority_class = class_counts.idxmax()
minority_class = class_counts.idxmin()

# Separate each class
df_majority = df_combined[df_combined['Risk Rating'] == majority_class]
df_minority = df_combined[df_combined['Risk Rating'] == minority_class]

# Undersample the majority class to 4500 records
df_majority_under = resample(df_majority, replace=False, n_samples=4500, random_state=42)

# Oversample the minority class to 4500 records if needed
df_minority_over = resample(df_minority, replace=True, n_samples=4500, random_state=42)
```

```
# Combine the resampled classes
df_resampled = pd.concat([df_majority_under, df_minority_over])

# Separate X and y again after resampling
X_resampled = df_resampled.drop(columns=['Risk Rating'])
y_resampled = df_resampled['Risk Rating']

# Check class distribution after resampling
print(f"Class distribution after resampling: {y_resampled.value_counts()}")

# Split the dataset into training and testing sets
X_resampled_train, X_resampled_test, y_resampled_train, y_resampled_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Class distribution after resampling: Risk Rating

| | |
|---|------|
| 1 | 4500 |
| 0 | 4500 |

Name: count, dtype: int64

```
[52]: print(f"Training samples: {X_resampled_train.shape}")
print(f"Validation samples: {X_resampled_test.shape}")
print(f"Training samples: {y_resampled_train.shape}")
print(f"Validation samples: {y_resampled_test.shape}")

Training samples: (8400, 18)
Validation samples: (2100, 18)
Training samples: (8400,)
Validation samples: (2100,)
```

Model selection

Model training is the process of teaching a machine learning model to identify patterns and make predictions using data. This involves feeding the model a large dataset consisting of input examples and their corresponding outcomes. Through this data, the model learns to recognize the relationships between the inputs and the target outputs by making iterative adjustments to its internal parameters, allowing it to improve its accuracy over time.

Our Model are,

- Logistic Regression Model
- K - Neighbors Classifier
- Support Vector Machine
- Random Forest Classifier
- Gradient Boosting Classifier


```
[60]: from sklearn.linear_model import LogisticRegression
      from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
      from sklearn.neighbors import KNeighborsClassifier
      from sklearn.svm import SVC
      from sklearn.metrics import accuracy_score, classification_report
      from sklearn.model_selection import cross_val_score

      # List of models to evaluate
      models = {
          "Logistic Regression": LogisticRegression(multi_class='multinomial', solver='lbfgs', max_iter=500),
          "Random Forest": RandomForestClassifier(),
          "Gradient Boosting": GradientBoostingClassifier(),
          "K-Nearest Neighbors": KNeighborsClassifier(),
          "SVM": SVC(kernel='linear')
      }

      # Evaluate each model using cross-validation
      for name, model in models.items():
          model.fit(X_train, y_train)
          y_pred = model.predict(X_test)
          accuracy = accuracy_score(y_test, y_pred)
          print(f"{name}: Accuracy = {accuracy}")
          print(classification_report(y_test, y_pred))
```

1. Logistic Regression Model

Linear classification; mostly applied to binary classification problems. It models probabilities with a logistic function and predicts class labels by selecting probabilities that are maximum. In the code, the algorithm has been set up for binary-class classification.

```
Logistic Regression: Accuracy = 0.9576190476190476
precision    recall  f1-score   support

      0       0.83      0.89      0.86       300
      1       0.98      0.97      0.98      1800

 accuracy          0.96       2100
 macro avg       0.90      0.93      0.92       2100
weighted avg       0.96      0.96      0.96       2100
```

2. Random Forest Classifier

The ensemble learning method works on the construction of multiple decision trees and then combines the results for a better classification outcome. It helps in reducing overfitting and improves generalization.

```
Random Forest: Accuracy = 0.9571428571428572
precision    recall  f1-score   support

      0       0.83      0.89      0.86       300
      1       0.98      0.97      0.97      1800

 accuracy          0.96       2100
 macro avg       0.90      0.93      0.92       2100
weighted avg       0.96      0.96      0.96       2100
```

3. Gradient Boosting Classifier

Another ensemble technique in which models are built in sequence, with each correcting the errors of the previously built model. It uses boosting for improving performance from weaker models, hence powerful for complicated datasets.

```

Gradient Boosting: Accuracy = 0.9566666666666667
      precision    recall  f1-score   support

      0       0.82       0.89       0.85        300
      1       0.98       0.97       0.97       1800

 accuracy          0.96          2100
  macro avg       0.90       0.93       0.91       2100
 weighted avg     0.96       0.96       0.96       2100

```

4. K-Nearest Neighbors Classifier

A simple instance-based learning algorithm that classifies an object based on the majority vote of its neighbors. It is easy to understand but computationally expensive for large data sets.

```

K-Nearest Neighbors: Accuracy = 0.9352380952380952
      precision    recall  f1-score   support

      0       0.82       0.71       0.76        300
      1       0.95       0.97       0.96       1800

 accuracy          0.94          2100
  macro avg       0.88       0.84       0.86       2100
 weighted avg     0.93       0.94       0.93       2100

```

5. Support Vector Machine (SVMs)

This is a classification algorithm that works by essentially finding the best hyperplane that separates the classes. The keyword `kernel='linear'` simply specifies a linear decision boundary, and is working best for data that are linearly separable.

```

SVM: Accuracy = 0.9576190476190476
      precision    recall  f1-score   support

      0       0.83       0.89       0.86        300
      1       0.98       0.97       0.98       1800

 accuracy          0.96          2100
  macro avg       0.90       0.93       0.92       2100
 weighted avg     0.96       0.96       0.96       2100

```

Model training

```
[65]: import joblib

# Assuming RandomForest was the best-performing model
best_model = RandomForestClassifier()
best_model.fit(X_train, y_train)

# Save the model to a file
joblib.dump(best_model, 'best_model.pkl')
print("Model saved successfully!")
```

Model saved successfully!

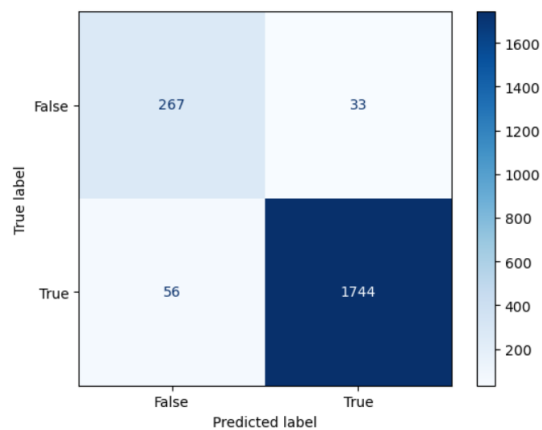
Confusion Matrix for Model Performance Evaluation

```
[39]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
import matplotlib.pyplot as plt

# Calculate the confusion matrix
cm = confusion_matrix(y_test, y_pred)

# Create a confusion matrix display with custom labels ('False', 'True')
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=['False', 'True'])

# Plot the confusion matrix with the 'Blues' color map
disp.plot(cmap='Blues')
plt.show()
```



Hyperparameter tuning

```
[67]: from sklearn.model_selection import GridSearchCV

# Initialize RandomForestClassifier
rf = RandomForestClassifier(random_state=42)

# Define hyperparameters to tune
param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

# Perform GridSearchCV to find the best hyperparameters
grid_search = GridSearchCV(estimator=rf, param_grid=param_grid, cv=5, scoring='accuracy', n_jobs=-1, verbose=2)

# Train the model
grid_search.fit(X_train, y_train)

# Get the best parameters
print(f"Best parameters: {grid_search.best_params_}")

# Evaluate on test data
y_pred = grid_search.predict(X_test)
print(f"Test Accuracy: {accuracy_score(y_test, y_pred)}")
print(classification_report(y_test, y_pred))
```

Fitting 5 folds for each of 81 candidates, totalling 405 fits
Best parameters: {'max_depth': 20, 'min_samples_leaf': 2, 'min_samples_split': 5, 'n_estimators': 100}
Test Accuracy: 0.9576190476190476

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.83 | 0.89 | 0.86 | 300 |
| 1 | 0.98 | 0.97 | 0.98 | 1800 |
| accuracy | | | 0.96 | 2100 |
| macro avg | 0.90 | 0.93 | 0.92 | 2100 |
| weighted avg | 0.96 | 0.96 | 0.96 | 2100 |

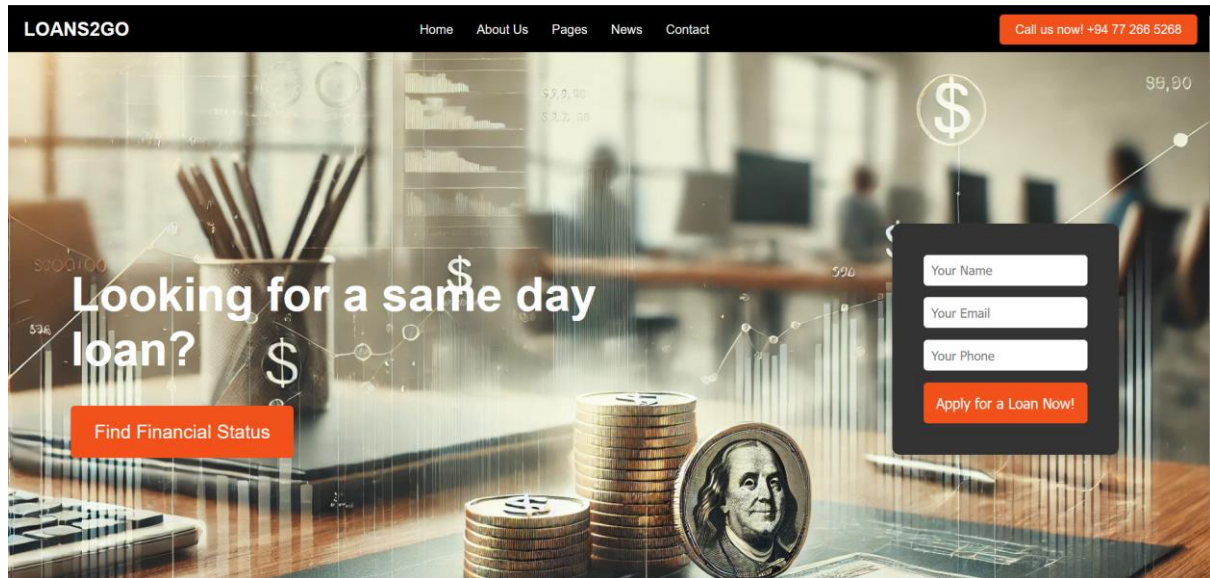
Cross-validation

```
[69]: # Cross-validation on training data
cv_scores = cross_val_score(grid_search.best_estimator_, X_train, y_train, cv=5, scoring='accuracy')

print(f"Cross-validation accuracy: {cv_scores.mean()}")

Cross-validation accuracy: 0.9528571428571428
```

Frontend



Risk Rating Prediction

Gender

Male

Age

1

0

100

City

Port Elizabeth

Payment History

Bad

Marital Status

Single

Loan Purpose

Auto

Loan Amount

10010

Income

50000

Employment Status

Employed

Years at Current Job

0

0

40

Predict Risk Rating

Conclusion

This project successfully proposes a machine learning-based solution for financial risk prediction that can enable financial institutions to make more appropriate assessments in the case of loan applications. Using demographic, financial, and behavioral data, we proposed the following challenges in the implementation: handling imbalanced classes via SMOTE, rectification of wrong and inconsistent data across columns, and removal of irrelevant features that would degrade the model's performance.

The final model optimized on cross-validation and hyperparameter tuning is therefore accurate in the prediction of risk levels such as Low or High. This system not only makes better decisions on loan approvals but also automates the process of risk evaluation involved, saving much manual effort and enhancing impartiality in assessments.

The work, in the near future, will be extended to real-time data for dynamic predictions; explainability will also be implemented in AI models, and the dataset will be increased with more financial variables. This can help upscale and refine the system further to enhance its usefulness for the risk management of the financial sector.

Work Distribution

| Student | Roles and Responsibilities |
|-----------------------|---|
| Bandara K M W G L A | <ul style="list-style-type: none"> • Data selection, preparation, preprocessing • Implementing and Training the model • Testing the model • Model Evaluation • Tuning and Optimizing the model • Model Deployment • Data visualizing • Implement user Interface |
| Jayasinghe J M H C | <ul style="list-style-type: none"> • Data selection, preparation, preprocessing • Implementing and Training the model • Testing the model • Model Evaluation • Tuning and Optimizing the model • Model Deployment • Data visualizing • Documentation |
| Abeyasinghe E A | <ul style="list-style-type: none"> • Data selection, preparation, preprocessing • Implementing and Training the model • Testing the model • Model Evaluation • Tuning and Optimizing the model • Model Deployment • Data visualizing • Documentation |
| Wicramasinghe D A T N | <ul style="list-style-type: none"> • Data selection, preparation, preprocessing • Implementing and Training the model • Testing the model • Model Evaluation • Tuning and Optimizing the model • Model Deployment • Data visualizing • Documentation |
| Mesandu W M S | <ul style="list-style-type: none"> • Data selection, preparation, preprocessing • Implementing and Training the model • Testing the model • Model Evaluation • Tuning and Optimizing the model • Model Deployment • Data visualizing • Documentation |